

Given a **query image patch** containing a fine-grained object (e.g., a specific bird species, a car model, or a brand logo), retrieve all images from a **large database** that contain the **same object**, even if the object appears at **different scales, partial occlusions**, or is embedded in **cluttered scenes**.

You are provided with:

- A **query image patch** (small cropped region with the object of interest).
- A **search set of images** (large, high-resolution, cluttered images with multiple objects).
- Bounding box annotations (for training only) for the object in some of the search images.

Your goal:

1. **Frame this problem formally as a deep learning task.**
2. **Propose an end-to-end CNN-based model** that can retrieve all images containing the same object (instance-level retrieval).
3. **Design an appropriate loss function** to handle **scale invariance, partial occlusion, and background clutter**.
4. Optionally, propose **efficiency improvements** for large-scale retrieval.

You need to submit:

- Report (max 6 pages) with the following sections:
 - Problem Formulation
 - Model Architecture
 - Loss Function
 - Evaluation Strategy
 - Conclusion